

Basics of Regression in Stata

LPO 9952 | Spring 2021

Intro

Stata was made for regression. It has the most advanced suite of regression functions and the easiest to use interface of any statistical programming environment. This session will get you started with how to estimate parameters for the simple regression model in STATA.

We'll be using data from the National Longitudinal Survey of Youth, 1997. For more information about the NLSY 97 sample, [click here](#).

Simple regression model

We'll be working with the same regression model as Wooldridge, with y as a linear function of x .

We're interested in coming up with estimates of the unknown population parameters.

Since we'll be doing OLS, we'll make all of the standard assumptions:

- The function is linear in parameters
- Our sample, including data y_i and x_i has been drawn randomly.
- There's variation in x
- The expected value of the error given the covariate is 0: $E(u|x)=0$, and the same is true in the sample, $E(ui|xi)=0$, meaning that x is fixed in repeated samples

The estimators

$\hat{\beta}_0$, $\hat{\beta}_1$ are unbiased given the above assumptions hold. This means that $E(\hat{\beta}_1 - \beta_1) = 0$ in repeated sampling.

Let's figure out how income and postsecondary attainment are related. Using the NLSY97 data set, we will get estimates for the following population regression model:

```
. version 15

. capture log close

. local gtype eps

. clear
```

```

. capture

. use nlsy97, clear
(Written by R.          )

. set seed 070328

. sample 10
(8,086 observations deleted)

. local y yinc

. local x ccol

. local ytitle "Income"

. local xtitle "Months of College"

```

Plotting Data

Before we do this, let's do a scatterplot. The scatterplot is the most fundamental graphical tool for regression. As a starting rule, never run a regression before looking at a scatterplot. In the accompanying do file, I've included the macros for setting this up in terms of x and y .

First, let's just plot y as a function of x :

```

. graph twoway scatter `y' `x', msize(small) ytitle(`ytitle') xtitle(`xtitle')

. graph export "simple_scatter.`gtype'", replace
(file simple_scatter.eps written in EPS format)

```

We can then add a lowess fit to see what the shape of the relationship between x and y looks like.

There are a variety of ways to check on the pattern on the data. A lowess regression gives you a local average estimate, which is sensitive to the patterns in the data:

```

. graph twoway lowess `y' `x', msize(small) ytitle(`ytitle') xtitle(`xtitle')

. graph export "simple_lowess.`gtype'", replace
(file simple_lowess.eps written in EPS format)

. graph twoway lowess `y' `x' || ///

```

```

scatter `y' `x', ///
msize(tiny) ///
msymbol(smcircle) ///
ytitle(`ytitle') ///
xtitle(`xtitle') ///
legend( order(2 "`xtitle'" 1 "Lowess fit") )

```

```

. graph export "scatter_lowess.gtype", replace
(file scatter_lowess.eps written in EPS format)

```

Our next step is to plot a linear fit to the data.

```

. graph twoway lfit `y' `x' || ///
scatter `y' `x', ///
msize(tiny) ///
msymbol(circle) ///
ytitle(`ytitle') ///
xtitle(`xtitle') ///
legend( order(2 `xtitle' 1 "Linear fit") ) //

```

Months not an integer, option order() ignored

```

. graph export "scatter_linear.gtype",replace
(file scatter_linear.eps written in EPS format)

```

Estimating Regression in Stata

We start with a basic regression of income on months of postsecondary education. There are a couple of ways of describing this, one is just to say we estimate a regression predicting income as a function of postseconarcy attendance. Another way is to sa y we regress income on postsecondary attendance.

```

. reg `y' `x'

```

Source	SS	df	MS	Number of obs	=	898
Model	6.5359e+10	1	6.5359e+10	F(1, 896)	=	110.03
Residual	5.3225e+11	896	594029160	Prob > F	=	0.0000
				R-squared	=	0.1094
				Adj R-squared	=	0.1084
Total	5.9761e+11	897	666230835	Root MSE	=	24373

yinc	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]

ccol	284.7932	27.1507	10.49	0.000	231.5068	338.0795
_cons	12366.25	1085.737	11.39	0.000	10235.36	14497.13

Quick Exercise

Run a regression with same dependent variable but a different independent variable. Interpret the results in one sentence. Write this sentence down.

Extracting Results

One key skill for today is being able to extract individual parts of the regression estimates from what Stata stores in memory. You need to build a map from the equations we'll be discussing to what can be accessed in Stata. Below, we start by extracting the regression coefficients.

```
. mat betamat=e(b)
```

The standard errors are stored as a variance-covariance matrix. To get a standard error, we need to take the square root of the elements of the diagonal of this matrix.

```
. mat vcmat=e(V)
```

```
. scalar myb=betamat[1,1]
```

```
. scalar varbeta1=vcmat[1,1]
```

```
. scalar sebeta1=sqrt(varbeta1)
```

We can use a different approach to get the same scalars. In Stata, referencing `_b[<varname>]` will pull the scalar for the coefficient associated with the variable name. Similarly, referencing `_se[<varname>]` will get the standard error for that coefficient.

```
. scalar beta0=_b[_cons]
```

```
. scalar li beta0
      beta0 = 12366.245
```

```
. scalar li
stat_sig = 9.044e-26
test = Significant
req_t = 1.9626151
my_pval = .05
myt = 10.829852
my_df = 896
adj_rsquare = .11476102
rsquare = .11574791
```

```

        fstat =    117.2857
        myf =    117.2857
    residss_std = 6.425e+08
    modss_std = 7.536e+10
        df_m =      1
    df_resid =      896
    my_mss = 7.536e+10
    ybar = 21186.17
    modss = 7.536e+10
    my_rss = 5.757e+11
    residss = 5.757e+11
        myk =      2
        myN =      898
    se_beta1 = 27.998234
    beta1 = 303.21674
    se_beta0 = 1145.8177
    beta0 = 12366.245
    sebeta1 = 27.150696
    varbeta1 = 737.16031
        myb = 284.79317

. scalar se_beta0=_se[_cons]

. scalar beta1=_b[`x']

. scalar se_beta1=_se[`x']

. scalar li beta1
    beta1 = 284.79317

```

Quick exercise: using both of the above methods, extract the estimate for the intercept

Confidence Intervals

By default, Stata gives 95% confidence intervals. To get confidence intervals at a different level, use the following code:

```
. reg `y' `x', level(90)
```

Source	SS	df	MS	Number of obs	=	898
Model	6.5359e+10	1	6.5359e+10	F(1, 896)	=	110.03
Residual	5.3225e+11	896	594029160	Prob > F	=	0.0000
				R-squared	=	0.1094
				Adj R-squared	=	0.1084
Total	5.9761e+11	897	666230835	Root MSE	=	24373

	yinc	Coef.	Std. Err.	t	P> t	[90% Conf. Interval]	
	ccol	284.7932	27.1507	10.49	0.000	240.088	329.4983
	_cons	12366.25	1085.737	11.39	0.000	10578.52	14153.97

Quick exercise: run the regression again, but this time get 80% CI

Residuals and Predictions

Residuals are not stored as part of the estimation results, but can be generated through the `predict` command. Below, I use the `predict` command to get residuals for this estimation.

```
. predict uhat, residuals

. tabstat uhat, stat(sum)

      variable |          sum
-----+-----
      uhat |    .0166245
-----+-----
```

These residuals can then be plotted as a function of x .

```
. graph twoway scatter uhat `x',yline(0) msize(tiny)

. graph export "residplot.`gtype'",replace
(file residplot.eps written in EPS format)

. graph twoway scatter uhat `x', ///
      msize(tiny) ///
      msymbol(circle) ///
      || ///
      scatter `y' `x', ///
      msize(tiny) ///
      msymbol(triangle)

. graph twoway scatter uhat `x', ///
      msize(tiny) ///
      msymbol(circle) ///
      || ///
      scatter `y' `x', ///
      msize(tiny) ///
```

```

        msymbol(triangle) ///
        || ///
    lfit `y' `x', ///
    lwidth(thin) ///
    yline(0, lpattern(dash) lwidth(thin)) ///
    legend(order(1 2 "Actual `ytitle'" 3))

```

```

. graph export "residplot_fancy.`gtype'",replace
(file residplot_fancy.eps written in EPS format)

```

The predicted value of y is also generated via the `predict` command.

```

. predict yhat
(option xb assumed; fitted values)

. graph twoway scatter yhat `x', ///
    msize(tiny) ///
    msymbol(circle) ///
    || ///
    scatter `y' `x', ///
    msize(tiny) ///
    msymbol(triangle)

```

```

. graph export "predict.`gtype'",replace
(file predict.eps written in EPS format)

```

These predicted values can be plotted relative to the actual data.

Measures of Model Fit

The first measure of model fit we consider is the F statistic. There are several ways to think about the F statistic. For now, I'm going to suggest that you think of it as the ratio of two measures. The first measure is the difference between the predicted value and the mean, or how different are your predictions than what would be predicted using the unconditional mean. The second measure is the difference between the predicted value and the actual value. We'll discuss this in class, but you should have an intuitive sense as to why the former should be large relative to the latter.

```

. ereturn list

```

```

scalars:

```

```

    e(rank) = 2
    e(ll_0) = -10396.10542311828

```

```

e(l1) = -10344.10078390533
e(r2_a) = .1083733615927051
e(rss) = 532250127517.0856
e(mss) = 65358931841.07336
e(rmse) = 24372.71343480899
e(r2) = .1093673712230364
e(F) = 110.0264704543859
e(df_r) = 896
e(df_m) = 1
e(N) = 898

```

macros:

```

e(cmdline) : "regress yinc ccol, level(90)"
e(title) : "Linear regression"
e(marginsok) : "XB default"
e(vce) : "ols"
e(depvar) : "yinc"
e(cmd) : "regress"
e(properties) : "b V"
e(predict) : "regres_p"
e(model) : "ols"
e(estat_cmd) : "regress_estat"

```

matrices:

```

e(b) : 1 x 2
e(V) : 2 x 2

```

functions:

```

e(sample)

```

Below, I conduct a test of statistical significance “by hand” to show how this is done in Stata.

```

. scalar myN=e(N)

. scalar myk=colsof(betamat)


. scalar residss=e(rss)

. gen diff=`y'-yhat

. gen diff_sq=diff*diff

. tabstat diff_sq,stat(sum) save

```


variable	sum
diff_sq	5.32e+11

```
. mat mymat=r(StatTotal)
```

```
. scalar my_rss=mymat[1,1]
```

```
. scalar li residss my_rss
  residss = 5.323e+11
  my_rss = 5.323e+11
```

```
. scalar modss=e(mss)
```

```
. tabstat `y', stat(mean) save
```

variable	mean
yinc	19910.73

```
. mat mymat=r(StatTotal)
```

```
. scalar ybar=mymat[1,1]
```

```
. gen diff2=yhat-ybar
```

```
. gen diff2_sq=diff2*diff2
```

```
. tabstat diff2_sq, stat(sum) save
```

variable	sum
diff2_sq	6.54e+10

```
. mat mymat=r(StatTotal)
```

```
. scalar my_mss=mymat[1,1]
```

```
. scalar li modss my_mss
  modss = 6.536e+10
  my_mss = 6.536e+10
```

```

. scalar df_resid=myN-myk

. scalar df_m=myk-1

. scalar modss_std=my_mss/df_m

. scalar residss_std=my_rss/df_resid

. scalar myf=modss_std/residss_std

. scalar fstat=e(F)

. scalar li myf fstat
      myf = 110.02647
      fstat = 110.02647

. corr yhat `y'
(obs=898)

-----+-----
      |      yhat      yinc
yhat |      1.0000
yinc |      0.3307      1.0000

. scalar rsquare= e(r2)

. scalar adj_rsquare= 1-((1-rsquare)*((myN-1)/(myN-myk)))

. scalar my_df=myN-myk

. scalar myt=beta1/sebeta1

. scalar my_pval=.05

. scalar req_t=invttail(my_df,(my_pval/2))

. scalar test=cond(abs(myt)>=req_t,"Significant","Not significant")

. scalar stat_sig=(2*ttail(my_df,myt))

. exit

```