

Regression Models for Binary Outcomes

There are many times when the dependent variables do not come in the nice, continuous, normally distributed form that we assume in introductory regression courses. The real world, being less generous, provides data in all kinds of forms. The first thing to think about is the underlying *Data Generation Mechanism*. What provides the outcome you see? If you reasonably think that the data should be a continuous variable with a more or less normal distribution, you can use a variety of transformations to get it where you need it to be. However, if the underlying DGM is decidedly non-normal, you need to look into alternative models.

A quite common occurrence in education research is a DGM that provides binary data. A partial list includes:

- High School/ College Graduation: yes or no?
- College Entry: yes or no?
- Departure from teaching: yes or no?
- Failure to make state standards: yes or no?
- Employment after graduation: yes or no?

All of these share a common DGM that gives a binary response. We'll consider two models to deal with this kind of outcome variable: the outmoded (yet making a comeback) linear probability model, and logistic regression.

Linear Probability Model

In this class we'll consider a model which predicts the binary outcome of college attendance as a function of ses, demographics, and test scores. To start, we could estimate this using good old OLS and see what happens. This is known as a *linear probability model*.

$$y_i = \beta_0 + \mathbf{x}_i\beta + e_i \quad , y \in 0, 1$$

To do this in Stata, we would run the following:

```

. /* Linear Probability Model */
.
. reg `y' `ses' `demog' `tests'

```

Source	SS	df	MS	Number of obs =	13284
Model	496.008461	9	55.1120512	F(9, 13274) =	374.14
Residual	1955.28354	13274	.147301759	Prob > F =	0.0000
				R-squared =	0.2023
				Adj R-squared =	0.2018
Total	2451.29201	13283	.184543552	Root MSE =	.3838

f2evratt	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
byses1	.1230284	.0051494	23.89	0.000	.1129347 .133122
amind	-.0799666	.0375121	-2.13	0.033	-.1534955 -.0064376
asian	.1061966	.0118984	8.93	0.000	.0828741 .1295192
black	.0664432	.0107626	6.17	0.000	.0453469 .0875395
hispanic	.0224763	.010479	2.14	0.032	.0019359 .0430167
multiracial	-.045084	.0158417	-2.85	0.004	-.0761361 -.014032
bysex	.0863008	.0067759	12.74	0.000	.0730191 .0995824
bynels2m	.0069647	.0003931	17.72	0.000	.0061943 .0077352
bynels2r	.0047414	.0005583	8.49	0.000	.003647 .0058357
_cons	.1329982	.0183363	7.25	0.000	.0970564 .16894

We'll use these coefficients as a baseline to compare with the logistic regressions we'll run later.

There are several problems with the LPM:

Probabilities outside of 0,1

In an ironic twist, the linear probability model does not give back probabilities. Instead, it fits a linear model to the data and gives predictions based on this assumption. This means that there can be predicted probabilities outside of 0,1. Of course, you can always just recode these predictions to be 1 or 0 but we generally try to avoid making up data.

Quick Exercise

Run a regression that models attendance as a function of test scores. See if there are any predicted probabilities that are outside of the range 0,1.

Non-Normality of Errors

The results from the LPM produce residuals that are far from normal, they are in fact binomial. This follows from

$$y_i = 1 \Rightarrow e_i = 1 - (\beta_0 + \mathbf{x}_i\beta)$$

$$y_i = 0 \Rightarrow e_i = -(\beta_0 + \mathbf{x}_i\beta)$$




Figure 1: Predicted Probability of Attendance by SES, LPM

Quick Exercise

Generate the residuals from the above regression and plot them against math test scores.

Heteroscedasticity

Since the residuals are a function of y , this means that the residuals are also a function of the systematic component of y represented by x . The variance of a binary variable is $p(1 - p)$, so the variance of e is:

$$V(e) = (\beta_0 + \mathbf{x}_i\beta)[1 - ((\beta_0 + \mathbf{x}_i\beta))] \quad (1)$$

When residuals are dependent on the data, we have heteroscedasticity.

Quick Exercise

Plot the residuals against the test scores variable. What do you see?

Of course, all of the normal methods for dealing with heteroscedasticity work here. Instead of the above, we can run ordinary least squares with robust standard errors, which will deal with the non-iid terms adequately.

```
. reg `y' `ses' `demog' `tests', robust
```

Linear regression

Number of obs = 13284
F(9, 13274) = 408.19
Prob > F = 0.0000
R-squared = 0.2023
Root MSE = .3838

	Coef.	Robust Std. Err.	t	P> t	[95% Conf. Interval]	
f2evratt						
byses1	.1230284	.0051319	23.97	0.000	.1129691	.1330876
amind	-.0799666	.0453305	-1.76	0.078	-.1688208	.0088877
asian	.1061966	.0105472	10.07	0.000	.0855226	.1268707
black	.0664432	.0117838	5.64	0.000	.0433454	.0895411
hispanic	.0224763	.01161	1.94	0.053	-.000281	.0452335
multiracial	-.045084	.0166198	-2.71	0.007	-.0776612	-.0125069
bysex	.0863008	.0067757	12.74	0.000	.0730194	.0995821
bynels2m	.0069647	.0004003	17.40	0.000	.00618	.0077495
bynels2r	.0047414	.000568	8.35	0.000	.003628	.0058547
_cons	.1329982	.0195842	6.79	0.000	.0946104	.1713861

Models for Binary Dependent Variables

Given the above theoretical problems with the LPM, almost all analysts favor the use of models specifically for binary variables. In particular, if you'd like to generate predicted probabilities as opposed to just understanding the conditional expectation function, you should use a model for a binary variable.

A notable exception to analysts preferring a nonlinear model for a binary dependent variable would be Angrist and Pischke, who say “The upshot of this discussion is that while a nonlinear model may fit the CEF (conditional expectation function) for LDVS (limited dependent variables) more closely than a linear model, when it comes to marginal effects, this probably matters little. This optimistic conclusion is not a theorem, but . . . it seems to be fairly robustly true.”

Generalized Linear Models

A *Generalized Linear Model* posits that y is a function of the independent variables and the coefficients or other parameters via a *link function*:

$$P(y|\mathbf{x}) = G(\beta_0 + \mathbf{x}_1\beta)$$

In our case, we're interested in the probability that y is 1

$$P(y = 1|\mathbf{x}) = G(\beta_0 + \mathbf{x}_i\beta)$$

There are several functions that “map” onto a 0,1 continuum. The most commonly used is the logistic function, which gives us the *logit model*.

The logistic function is given by:

$$f(x) = \frac{1}{1 + \exp^{-k(x-x_0)}}$$

Mapped onto our GLM, this gives us:

$$P(y = 1|\mathbf{x}) = \frac{\exp(\beta_0 + \mathbf{x}_i\beta)}{1 + \exp(\beta_0 + \mathbf{x}_i\beta)}$$

The critical thing to note about the above is that the link function maps the entire result of estimation ($\beta_0 + \mathbf{x}_i\beta$) onto the 0,1 continuum. Thus, the change in the $P(y = 1|\mathbf{x})$ is a function of all of the independent variables and coefficients together, *not* one at a time.

What does this mean? It means that the coefficients can only be interpreted on the *logit* scale, and don’t have the normal interpretation we would use for OLS regression. Instead, to understand what the logistic regression coefficients mean, you’re going to have to convert the entire term ($\beta_0 + \mathbf{x}_i\beta$) to the probability scale, using the inverse of the function.

Maximum Likelihood Estimation

The other important thing to know about the logit model is that it is estimated via maximum likelihood. This is in contrast to OLS, which we estimate via direct computation. You could, given the time and the willingness, sit down and compute the estimates for an OLS model (I have! It’s . . . fun?). For a MLE model, we use a likelihood function which gives us the likelihood of the *data* given a certain set of proposed parameters. The computer then adjusts the estimates of the parameters using a certain algorithm, checks to see if the likelihood has gone up or down, and repeats. The algorithm keeps on looking for the parameters that make the data most likely. This stops at a pre-set point when improvements in the likelihood are small. This condition is called *convergence*

To run a logistic regression in Stata, enter:

```

. /*Logistic Regression */
.
. logit `y' `ses' `demog' `tests'

Iteration 0:   log likelihood = -7383.1405
Iteration 1:   log likelihood = -5998.4224
Iteration 2:   log likelihood = -5882.6198
Iteration 3:   log likelihood = -5881.8938
Iteration 4:   log likelihood = -5881.8937

Logistic regression               Number of obs   =       13284
                                LR chi2(9)           =       3002.49

```

```

Log likelihood = -5881.8937
Prob > chi2      = 0.0000
Pseudo R2       = 0.2033

```

f2evratt	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
byses1	.9044976	.0375964	24.06	0.000	.8308101	.9781852
amind	-.2907766	.2208319	-1.32	0.188	-.7235993	.142046
asian	1.024893	.0997533	10.27	0.000	.8293798	1.220406
black	.4880564	.0687739	7.10	0.000	.353262	.6228508
hispanic	.3010218	.0673856	4.47	0.000	.1689485	.4330951
multiracial	-.2907231	.1021569	-2.85	0.004	-.4909469	-.0904993
bysex	.5852038	.0470378	12.44	0.000	.4930114	.6773962
byncls2m	.0461185	.002718	16.97	0.000	.0407913	.0514457
byncls2r	.0308627	.0037971	8.13	0.000	.0234205	.0383049
_cons	-2.684406	.124637	-21.54	0.000	-2.92869	-2.440122

```

.
. est store full_model
.
. gen mysample=e(sample)

```

The iterations refer to the steps as Stata computes the log likelihood and then checks for convergence. You then get the final log likelihood, the number of observations, the χ^2 statistics for the likelihood ratio test, the p value for that χ^2 and the Pseudo R^2 . The coefficient table should look generally familiar.

Interpretation of Results

The coefficients from a logit model are not directly interpretable in the same way as OLS estimates. This is because they are linear only in the log odds, which is not a way that (ordinary) humans think about the world. We are interested in the changes in the $P(y = 1)$, which means that we have to take into account the *whole* model, since this is what is mapped onto the (0,1) scale.

Quick Exercise

Run a logistic regression predicting college attendance by math scores. Generate a predicted outcome, and plot this against math scores. Use the form `predict yhat, pr` to get predicted probabilities for $y = 1$.

Generating Marginal Effects

The marginal effect of a covariate is the predicted change in the probability of the outcome as a result of a one unit change in the covariates, with all other covariates held constant at some level. For logistic regression, it's most common to hold all other variables at their means, but it's not the only way to go.

In Stata, marginal effects can be computed quickly using the `margins` command. Running this for all variables with the `dydx` option will give you the marginal effect of each covariate,

interpreted as a change in probability of the outcome for a one unit change in the independent variable, holding all other variables constant.

```

. /* Generating marginal effects */
. margins, dydx(*) /*for all coefficients, default is to hold others at mean */

Average marginal effects          Number of obs   =       13284
Model VCE      : OIM

Expression   : Pr(f2evratt), predict()
dy/dx w.r.t. : byses1 amind asian black hispanic multiracial bysex bynels2m bynels2r

```

		Delta-method					
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]	
byses1		.1295129	.0050031	25.89	0.000	.1197071	.1393188
amind		-.0416356	.0316113	-1.32	0.188	-.1035926	.0203213
asian		.146752	.0141156	10.40	0.000	.1190861	.174418
black		.0698837	.0097859	7.14	0.000	.0507036	.0890637
hispanic		.0431026	.0096314	4.48	0.000	.0242253	.0619799
multiracial		-.041628	.0146135	-2.85	0.004	-.07027	-.012986
bysex		.083794	.006603	12.69	0.000	.0708524	.0967355
bynels2m		.0066036	.000375	17.61	0.000	.0058685	.0073387
bynels2r		.0044192	.0005392	8.20	0.000	.0033623	.005476

Generating Probabilities

Since we are primarily interested in the probability that $y = 1$, the best course to take to interpret coefficients is to demonstrate what happens to the dependent variable as we allow certain independent variables to move within their range. In Stata, the margins is indispensable for this task.

Reporting these results well takes some time and some effort. The analyst needs to think carefully about the likely changes in the independent variables, and the right levels at which to hold the other variables constant. For our first take, let's let race vary and then generate margins for a series of values of SES:

```

. /* Generate predicted probabilities over range of ses*/
.
. local x byses1

.
. sum `x', detail

      socio-economic status composite, v.1
-----
Percentiles      Smallest
1%             -1.54      -2.11
5%             -1.18      -2.11
10%            -.95       -1.97      Obs           15325
25%            -.5        -1.97      Sum of Wgt.    15325

50%            .04
              Largest
75%            .59         1.8
90%            1.05         1.8      Mean           .0426499
95%            1.26         1.81     Std. Dev.       .7429194
99%            1.55         1.82     Variance        .5519293
              Kurtosis      2.354114

```

```

.
. local no_steps=20

.
. local mymin=r(min)

. local mymax=r(max)

. local diff=`mymax`-`mymin`

. local step=`diff`/`no_steps`

.
.
. margins , predict(xb) ///
>   at((mean) _continuous ///
>       (min) `demog` ///
>       `x'=(`mymin`(`step`)`mymax`) ///
>       ) ///
>   post

Adjusted predictions                                Number of obs   =       13284
Model VCE      : Robust

Expression    : Linear prediction, predict(xb)

1._at      : byses1      =       -2.11
             amind       =           0 (min)
             asian       =           0 (min)
             black       =           0 (min)
             hispanic    =           0 (min)
             multiracial =           0 (min)
             bysex       =           1 (min)
             byncls2m    =      46.1288 (mean)
             byncls2r    =     30.23475 (mean)

```

We can plot the results from this margins command like so:

For another example, we'll let ses and race vary, holding other variables constant. The resulting graphic shows the predicted change in the probability of college attendance as these two variables change.

Quick Exercise

Generate a set of predicted probabilities from our model that show the predicted change in probability of college attendance for white women with high, medium and low test scores, holding other variables constant at a reasonable level. Do this for the entire range of ses and plot the result.

Reporting Odds Ratios

The odds ratio for any coefficient is the predicted change in odds for a one unit change in the covariate, while holding all other variables constant (usually at their mean). The simplest way to obtain odds ratios from Stata is to run the logit model with the `logistic` command:

To get the odds ratios, we calculate the exponent of the coefficients. The result show the ratio of the probability that $y = 1$ when we shift the covariate of interest by a specific amount.




Figure 2: Predicted Probability of Attendance by SES, Logit Model

In Stata, the default is one unit, but this can be adjusted. The interpretation of this is that after a change of the specified size the odds are so much larger or smaller.

```
.
.
. listcoef /*Display odds ratios from model in memory */

logit (N=13284): Factor Change in Odds

Odds of: Yes vs No
```

	f2evratt	b	z	P> z	e ^b	e ^b StdX	SDofX
byses1		0.90450	24.058	0.000	2.4707	1.9612	0.7447
amind		-0.29078	-1.317	0.188	0.7477	0.9743	0.0894
asian		1.02489	10.274	0.000	2.7868	1.3480	0.2913
black		0.48806	7.097	0.000	1.6291	1.1790	0.3373
hispanic		0.30102	4.467	0.000	1.3512	1.1098	0.3462
multiracial		-0.29072	-2.846	0.004	0.7477	0.9397	0.2138
bysex		0.58520	12.441	0.000	1.7954	1.3396	0.4997
bynels2m		0.04612	16.968	0.000	1.0472	1.8722	13.5976
bynels2r		0.03086	8.128	0.000	1.0313	1.3370	9.4115

Odds ratios for negative coefficients can be hard to interpret, so you can also ask Stata to

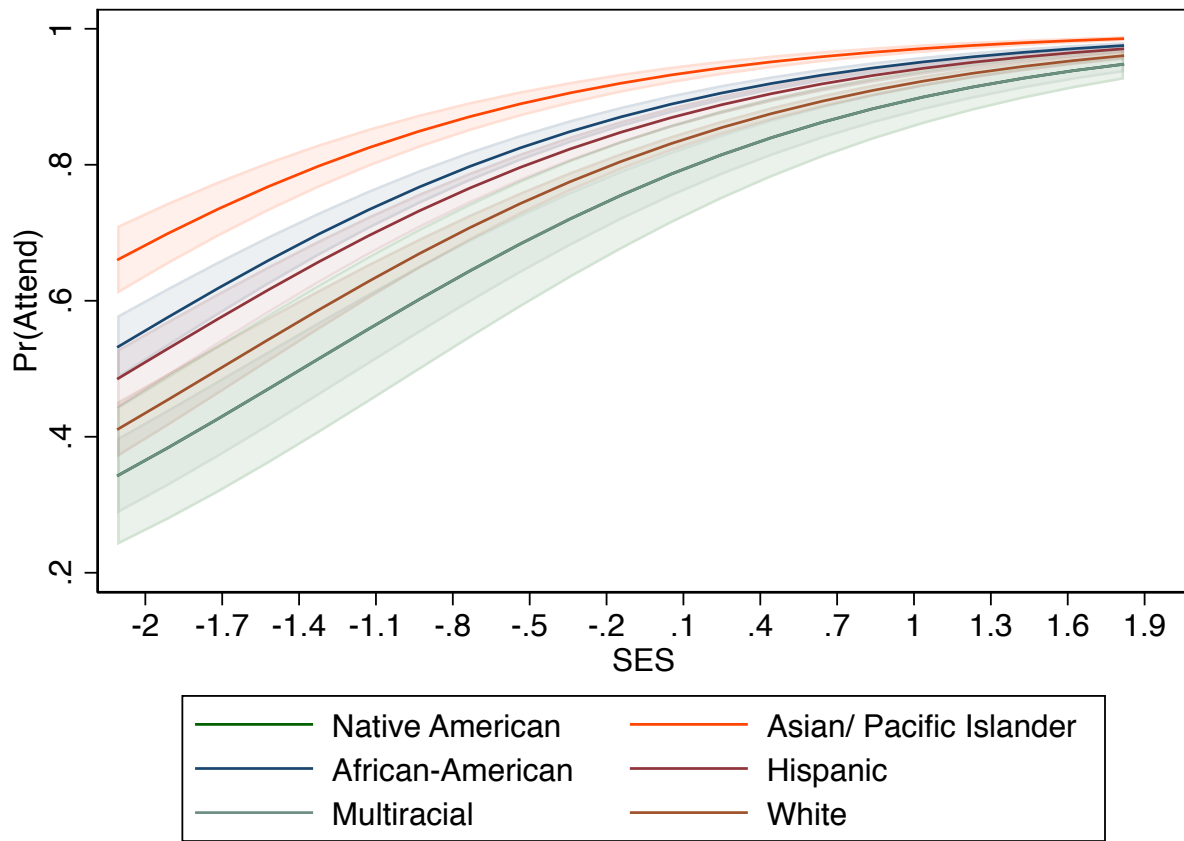


Figure 3: Predicted Probability of Attendance by SES and Race

reverse the interpretation, like so:

```
. listcoef, reverse /* Reverses interpretation, helps with negative coeffs */
logit (N=13284): Factor Change in Odds
Odds of: No vs Yes
```

	f2evratt	b	z	P> z	e ^b	e ^b StdX	SDofX
byses1		0.90450	24.058	0.000	0.4047	0.5099	0.7447
amind		-0.29078	-1.317	0.188	1.3375	1.0263	0.0894
asian		1.02489	10.274	0.000	0.3588	0.7419	0.2913
black		0.48806	7.097	0.000	0.6138	0.8482	0.3373
hispanic		0.30102	4.467	0.000	0.7401	0.9010	0.3462
multiracial		-0.29072	-2.846	0.004	1.3374	1.0641	0.2138
bysex		0.58520	12.441	0.000	0.5570	0.7465	0.4997
bynels2m		0.04612	16.968	0.000	0.9549	0.5341	13.5976
bynels2r		0.03086	8.128	0.000	0.9696	0.7479	9.4115

Odds ratios seem simple but are not really—they are a specific way of calculating predicted probabilities that many times allow the analyst to avoid thinking carefully about the results. Use

them with care. Or not at all. Really, you shouldn't use them. Nobody does a good job writing them up.

Quick Exercise

Interpret the odds ratio for income from the above regression. Write down your answer.

Hypothesis Testing

To test whether an individual coefficient is 0, we use a similar method as in OLS, except the distribution of the test statistic is the Normal distribution not the T distribution. For that reason, we report z and not t statistics.

Likelihood Ratio Test

To test whether a subset of coefficients jointly increase model fit, we use the likelihood ratio test. This is two times the difference in log likelihoods, and can be done quite easily in Stata. To test the full model in our example against one with just the kids variables, we would run:

In Stata, the test is:

```
.  
. lrtest full_model ses  
  
Likelihood-ratio test                                LR chi2(8) =   1398.05  
(Assumption: ses nested in full_model)              Prob > chi2 =    0.0000  
  
. quietly logit `y' `demog' if mysample==1  
  
. est store demog  
  
. lrtest full_model demog  
  
Likelihood-ratio test                                LR chi2(3) =   2542.87  
(Assumption: demog nested in full_model)            Prob > chi2 =    0.0000  
  
. quietly logit `y' `tests' if mysample==1  
  
. est store tests  
  
. lrtest full_model tests  
  
Likelihood-ratio test                                LR chi2(7) =    860.54  
(Assumption: tests nested in full_model)            Prob > chi2 =    0.0000
```

Goodness of Fit

Goodness of fit in this context is somewhat more complex than in OLS. There are three basic methods: the Chi Square Test, McFadden's Pseudo R^2 , and the percent correctly predicted method.

Chi Square Test

The Chi Square Test, like the F test, tests the null hypothesis that all of the coefficients are 0. Like the F test, it is a very weak test.

Pseudo R^2

The pseudo R^2 is just like it sounds—not a real R^2 . It shares the same general interpretation of R^2 —levels close to 0 are bad, levels close to one are good. It's based on transforming the likelihood ratio to be constrained by 0,1.

Percent Correctly Predicted

The last method to use is to think about the percent correctly predicted. This is fairly intuitive and easy to do. However, the key concept here is the threshold of correctness. Generally, we use .5 probability as the threshold, but there's nothing magic about .5. Depending on the application, we might be more comfortable with a lower false negative, or a lower false positive—it just depends.

The command in Stata is:

```
. /* Percent Correctly Predicted */
.
. estat classification

Logistic model for f2evratt

Classified |      True      |      Total
-----+-----+-----
      +   |      9322      |      11390
      -   |      719       |      1894
-----+-----+-----
    Total |     10041      |     13284

Classified + if predicted Pr(D) >= .5
True D defined as f2evratt != 0
-----+-----+-----
Sensitivity                Pr( +| D)    92.84%
Specificity                Pr( -| ~D)    36.23%
Positive predictive value  Pr( D| +)    81.84%
Negative predictive value  Pr(~D| -)    62.04%
-----+-----+-----
False + rate for true ~D   Pr( +| ~D)   63.77%
False - rate for true D    Pr( -| D)    7.16%
False + rate for classified + Pr(~D| +)   18.16%
False - rate for classified - Pr( D| -)   37.96%
```

Correctly classified	79.02%
----------------------	--------

The resulting table (also known as a confusion matrix) gives you both sensitivity (the proportion of 1's that are correctly identified) and specificity (the proportion of 0's correctly identified). Depending on the application, one or the other may be more important to the analyst.

ROC (an acronym for receiver-operator characteristic) goes some way towards overcoming the problem of choosing a threshold. The ROC considers what happens to the classification rate as the classification threshold ranges from zero to one. The ROC is based on changing the classification threshold. When doing this, the number of "true positive" classifications changes in inverse proportion to the number of "false positive" classifications. ROC analysis is usually done graphically, plotting the TPF (true positive fraction, $TPF = 1 - FNF$), versus the FPF (false positive fraction) as the classification threshold varies. The resulting function is defined on the unit square, and the area under the ROC curve c is interpreted as a measure of the classification success of the model. A value of $c = .5$ indicates random predictions and a value of $c = 1$ indicates perfect prediction.

The command in Stata is:

```
. /*Area under Receiver/Operator Characteristic Curve */
.
. lroc

Logistic model for f2evratt

number of observations =    13284
area under ROC curve   =    0.8013
```

You can also compare models and do a graphical comparison, like so:

```
. est restore ses
(results ses are active now)

.
. predict xb_ses, xb
(964 missing values generated)

.
. est restore full_model
(results full_model are active now)

.
. predict xb_full, xb
(964 missing values generated)

.
. roccomp f2evratt xb_full xb_ses, graph summary
```

	Obs	ROC Area	Std. Err.	-Asymptotic Normal-- [95% Conf. Interval]
xb_full	13284	0.8013	0.0043	0.79293 0.80960
xb_ses	13284	0.7300	0.0048	0.72056 0.73934

```
Ho: area(xb_full) = area(xb_ses)
chi2(1) =    345.06      Prob>chi2 =    0.0000
```

Figure 4: Area Under ROC Curve: Comparing Two Models

